


Characterizing Machine Learning I/O with MLPerf Storage

Oana Balmau

WoCC Keynote, December 12th, 2023



Humanity produces a lot of data



Source: IDC 2022

Humanity produces a lot of data




Source: IDC 2022

Data is the moving force of ML algorithms

... but in many projects the **storage decision is an afterthought**

Inefficient I/O can slow down ML Workloads

Dataset fits in system memory



Dataset = 2x system memory



Training time increased by 3x

Example: Image Segmentation with 3D U-Net

Medical image segmentation


2019 Kidney Tumor Segmentation Challenge (KiTS19)

CT scans from ~300 kidney tumor cases



An example of a coronal section of one of the training cases with its ground truth segmentation overlaid (kidney in red, tumor in blue).

Source <https://arxiv.org/pdf/1912.01054.pdf>



Sample images from the KiTS19 dataset before (left) and after (right) preprocessing.

Source: <https://arxiv.org/pdf/1908.02625.pdf>


Example: Image Segmentation with 3D U-Net

File System

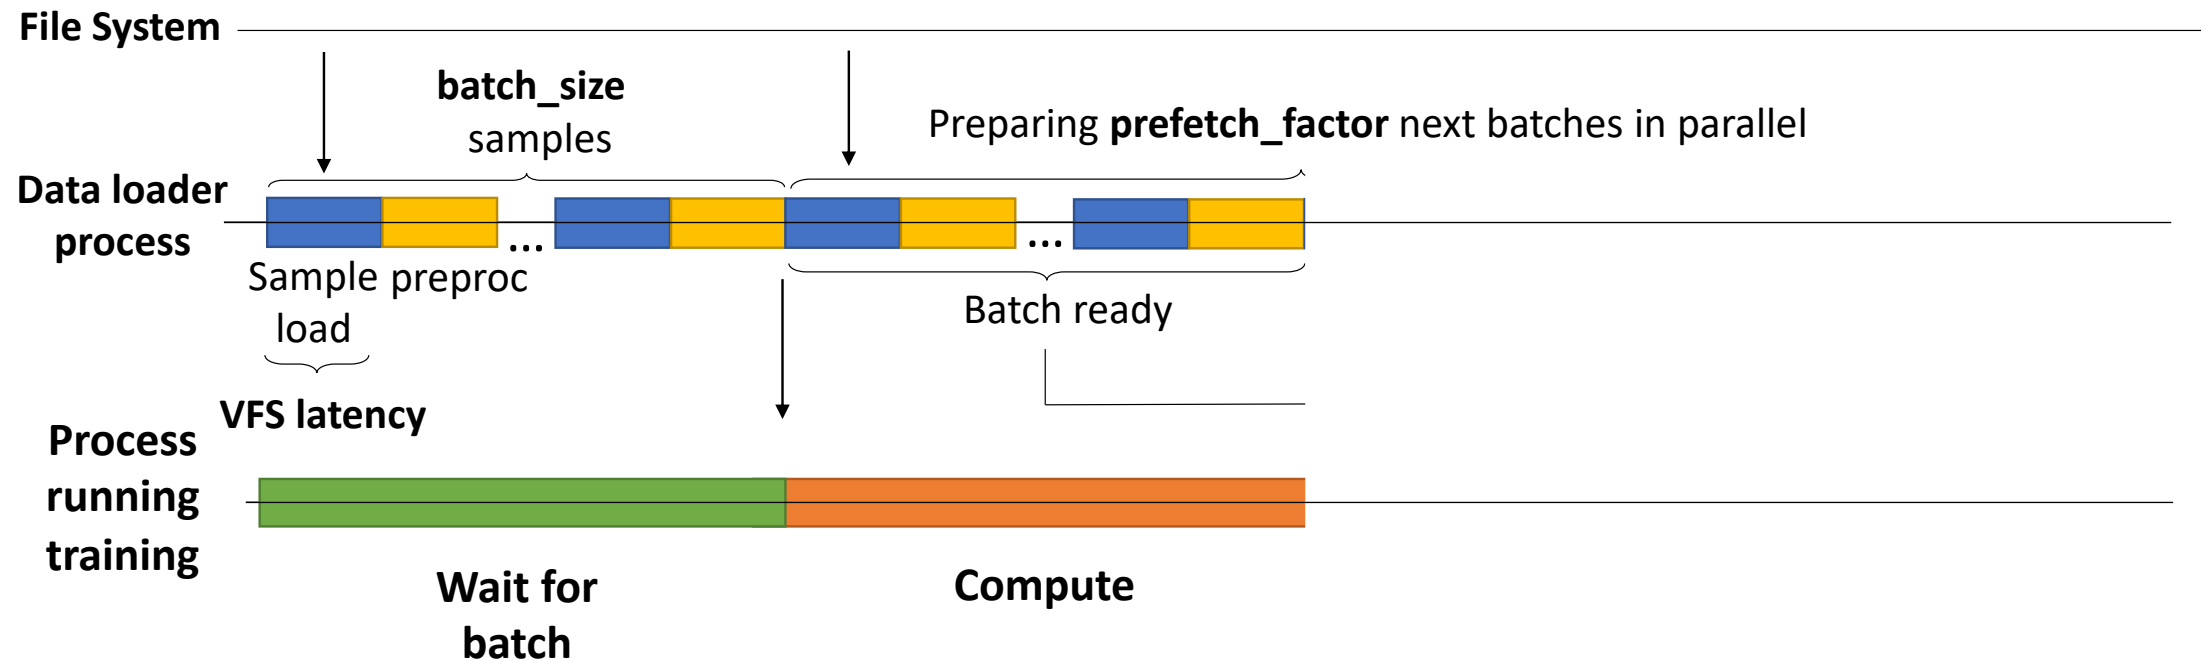
**Data loader
process**

**Process
running
training**

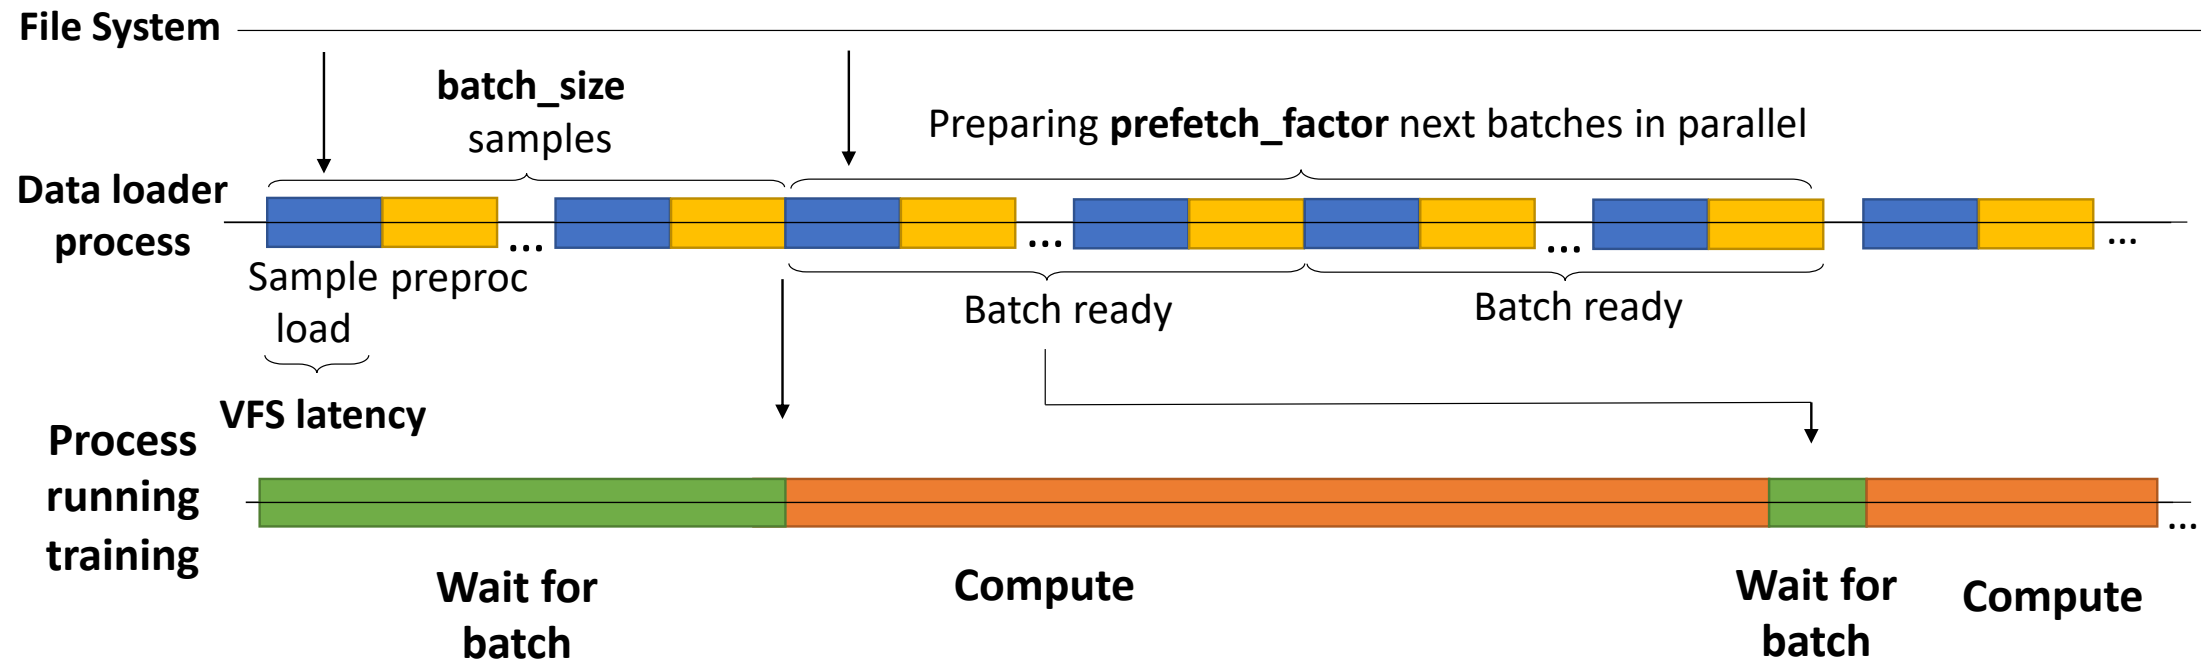
Example: Image Segmentation with 3D U-Net



Example: Image Segmentation with 3D U-Net



Example: Image Segmentation with 3D U-Net





Inefficient I/O can slow down ML Workloads

Experiment setup

- DGX-1 server
 - 8 x V100 GPUs, 32GB GPU memory
 - 512GB DRAM
- Image segmentation workload:
 - Unet3D, Pytorch
 - MLPerf Training implementation
 - KiTS19 dataset

Dataset fits in system memory



■ ML Training ■ ML Evaluation ■ Disk I/O Read ■ In-memory Read ■ GPU ■ CPU ■ GPU Memory

Inefficient I/O can slow down ML Workloads



Experiment setup

- DGX-1 server
 - 8 x V100 GPUs, 32GB GPU memory
 - 512GB DRAM
- Image segmentation workload:
 - Unet3D, Pytorch
 - MLPerf Training implementation
 - KiTS19 dataset

High GPU utilization →

Little disk access →

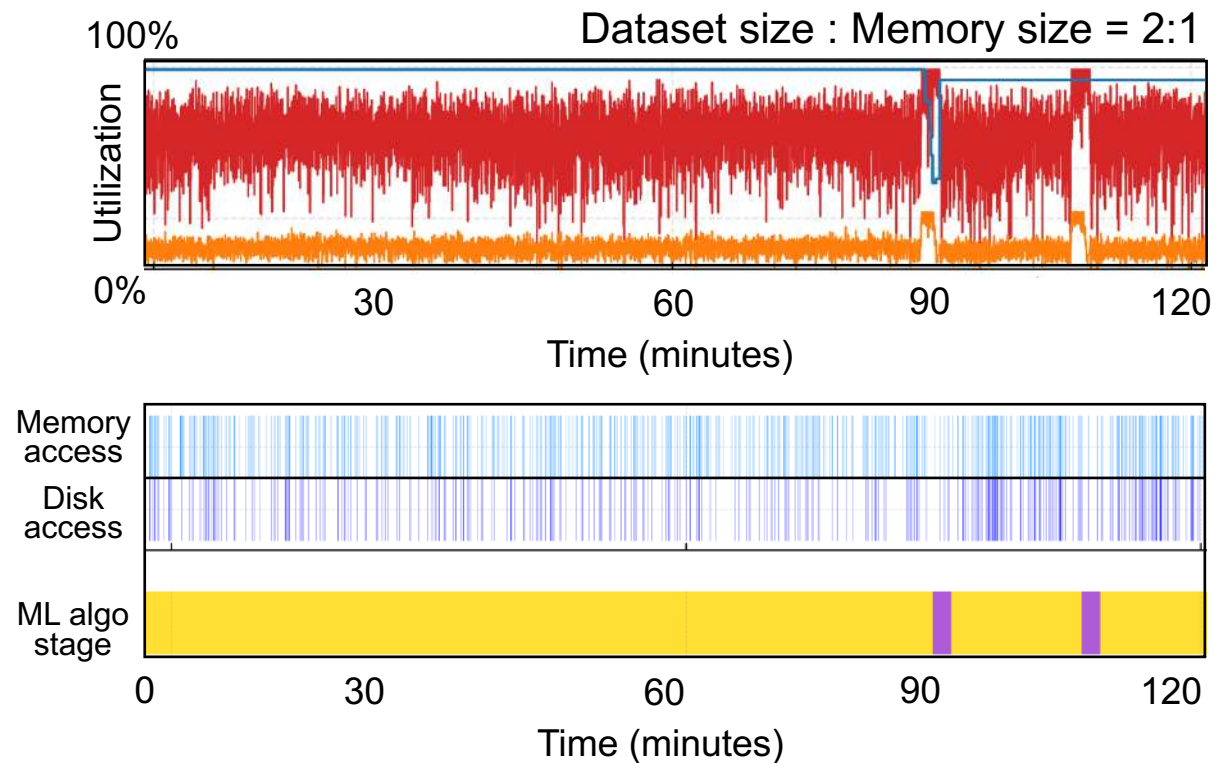
Dataset fits in system memory



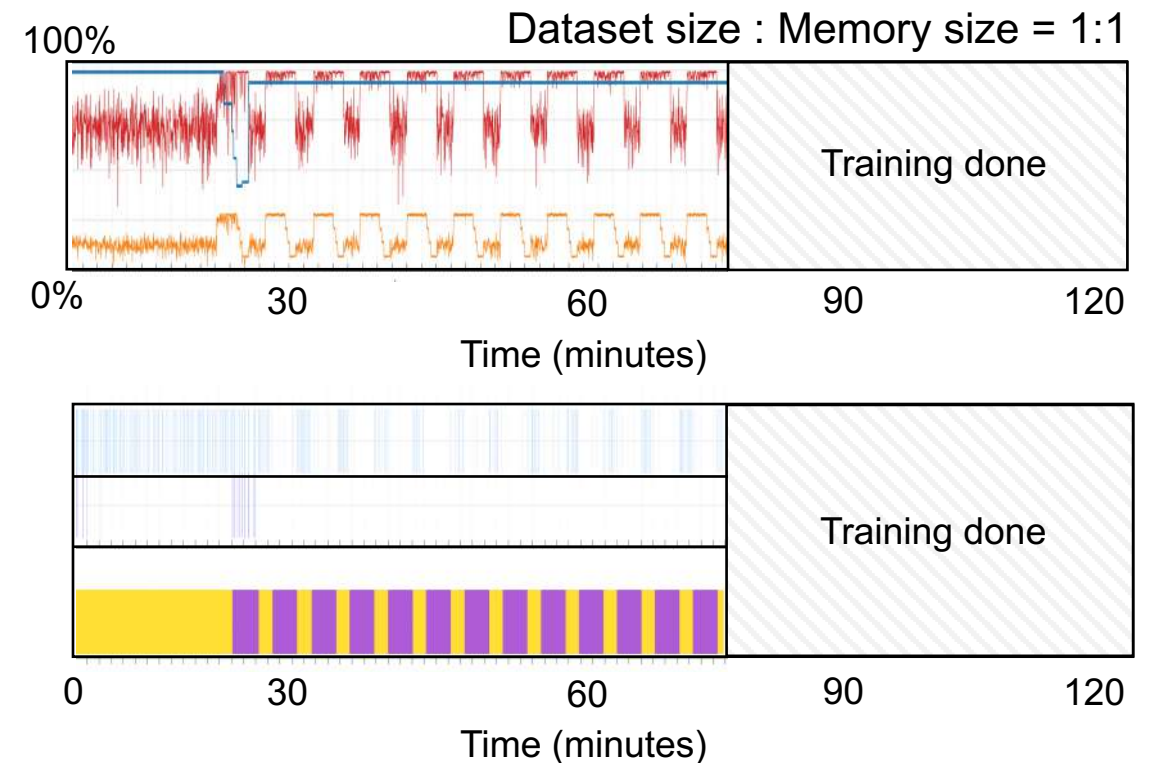
ML Training ML Evaluation Disk I/O Read In-memory Read GPU CPU GPU Memory

Inefficient I/O can slow down ML Workloads

Dataset does not fit in memory





Dataset fits in system memory





ML Training ML Evaluation Disk I/O Read In-memory Read GPU CPU GPU Memory

Inefficient I/O can slow down ML Workloads

Dataset does not fit in memory





Dataset fits in system memory





ML Training ML Evaluation Disk I/O Read In-memory Read GPU CPU GPU Memory

Inefficient I/O can slow down ML Workloads

Dataset does not fit in memory



Dataset fits in system memory



Training time increased by 3x

ML Training ML Evaluation Disk I/O Read In-memory Read GPU CPU GPU Memory

Data is the moving force of ML algorithms

... but in many projects the **storage decision is an afterthought**

Why create an ML Storage benchmark?

Why create an ML Storage benchmark?

- Understand storage bottlenecks in ML workloads
and propose optimizations
- Help AI/ML researchers and practitioners
make an informed storage decision

MLPerf Storage Working Group (132 members)

Who are we?



- Academia
- Storage Vendors
- Accelerator Vendors
- End Users



Current ML/AI benchmarks

Many existing ML/AI benchmarks



DeepMind Lab



MLPerf



OpenAI

DLBT



PMLDB




DAWNBench

Current ML/AI benchmarks

- Focus on **end-to-end testing**
 - hard to isolate value of each component
- Insist on **training and inference** speed
 - tend to simplify storage
 - ignore pre-processing
- **Expensive accelerators** needed to run
- Require **extensive entry knowledge**



PMLDB



Benchmark Vision

Existing benchmarks

Focus on **end-to-end testing**

Simplified storage setup

Expensive accelerators needed to run

Require **extensive entry knowledge**

Our work

Focus on **storage impact in ML/AI**

Realistic **storage & pre-processing** settings



No accelerator required to run

Minimal AI/ML knowledge required


Stages of the ML Pipeline



**Data cleaning &
pre-processing**

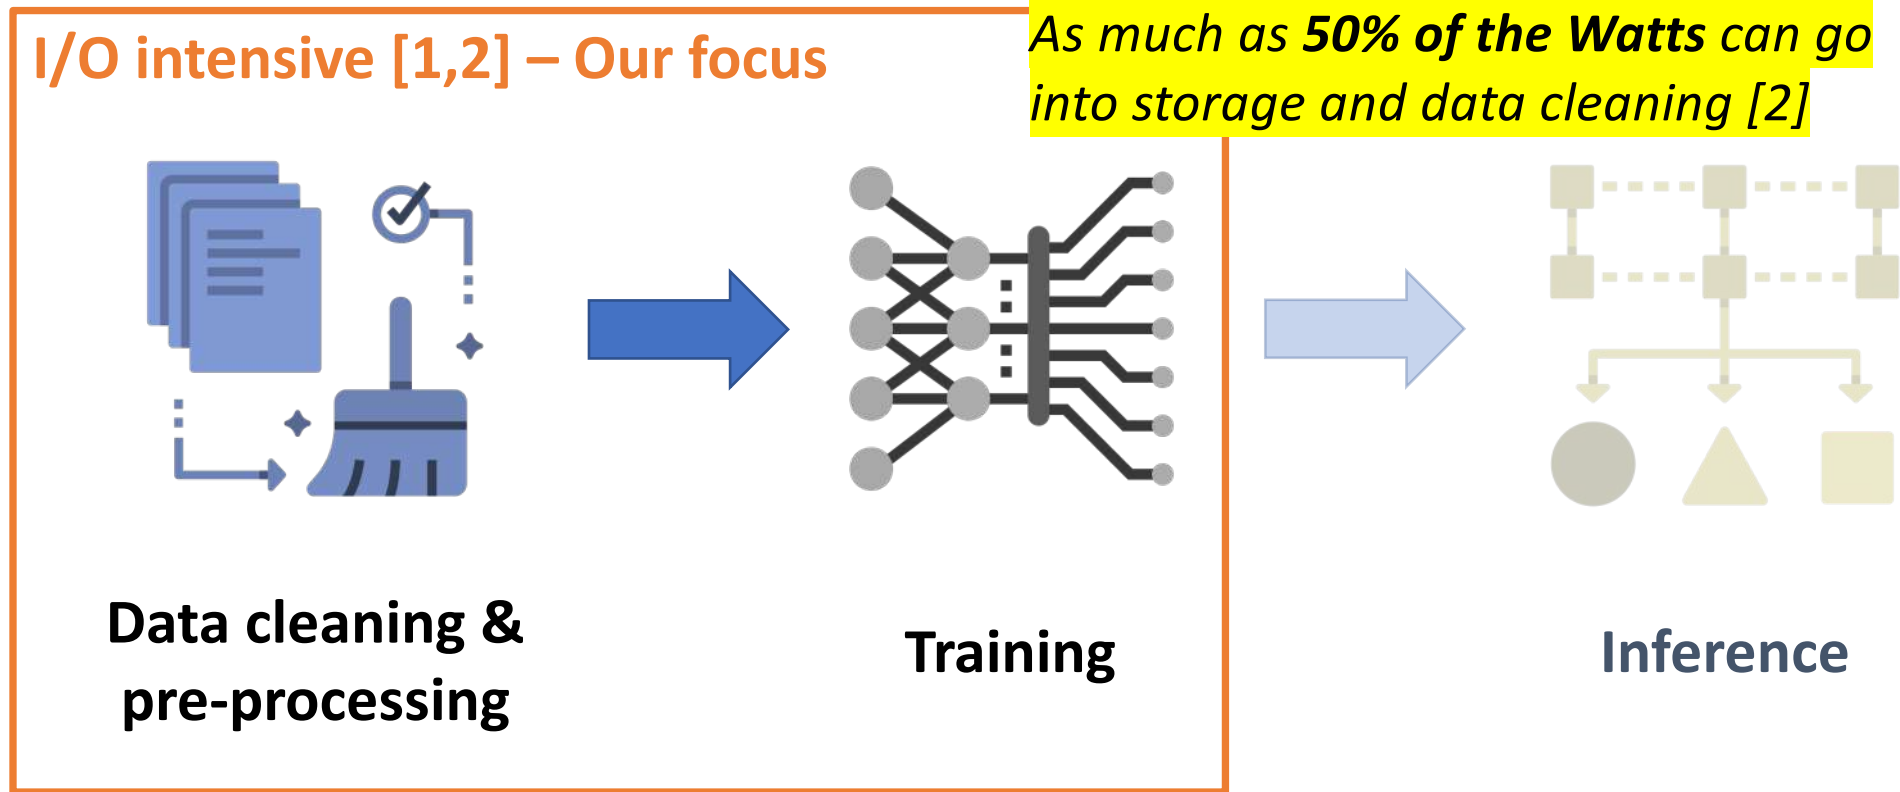


Training



Inference

Stages of the ML Pipeline



[1] Murray et al. *tf.data: A Machine Learning Data Processing Framework*, VLDB 21.

[2] Zhao et al. *Understanding Data Storage and Ingestion for Large-Scale Deep Recommendation Model Training* ISCA 22.



Data Pipeline in ML: Pre-processing

Storage resources

Disk



Memory

Compute resources

CPUs

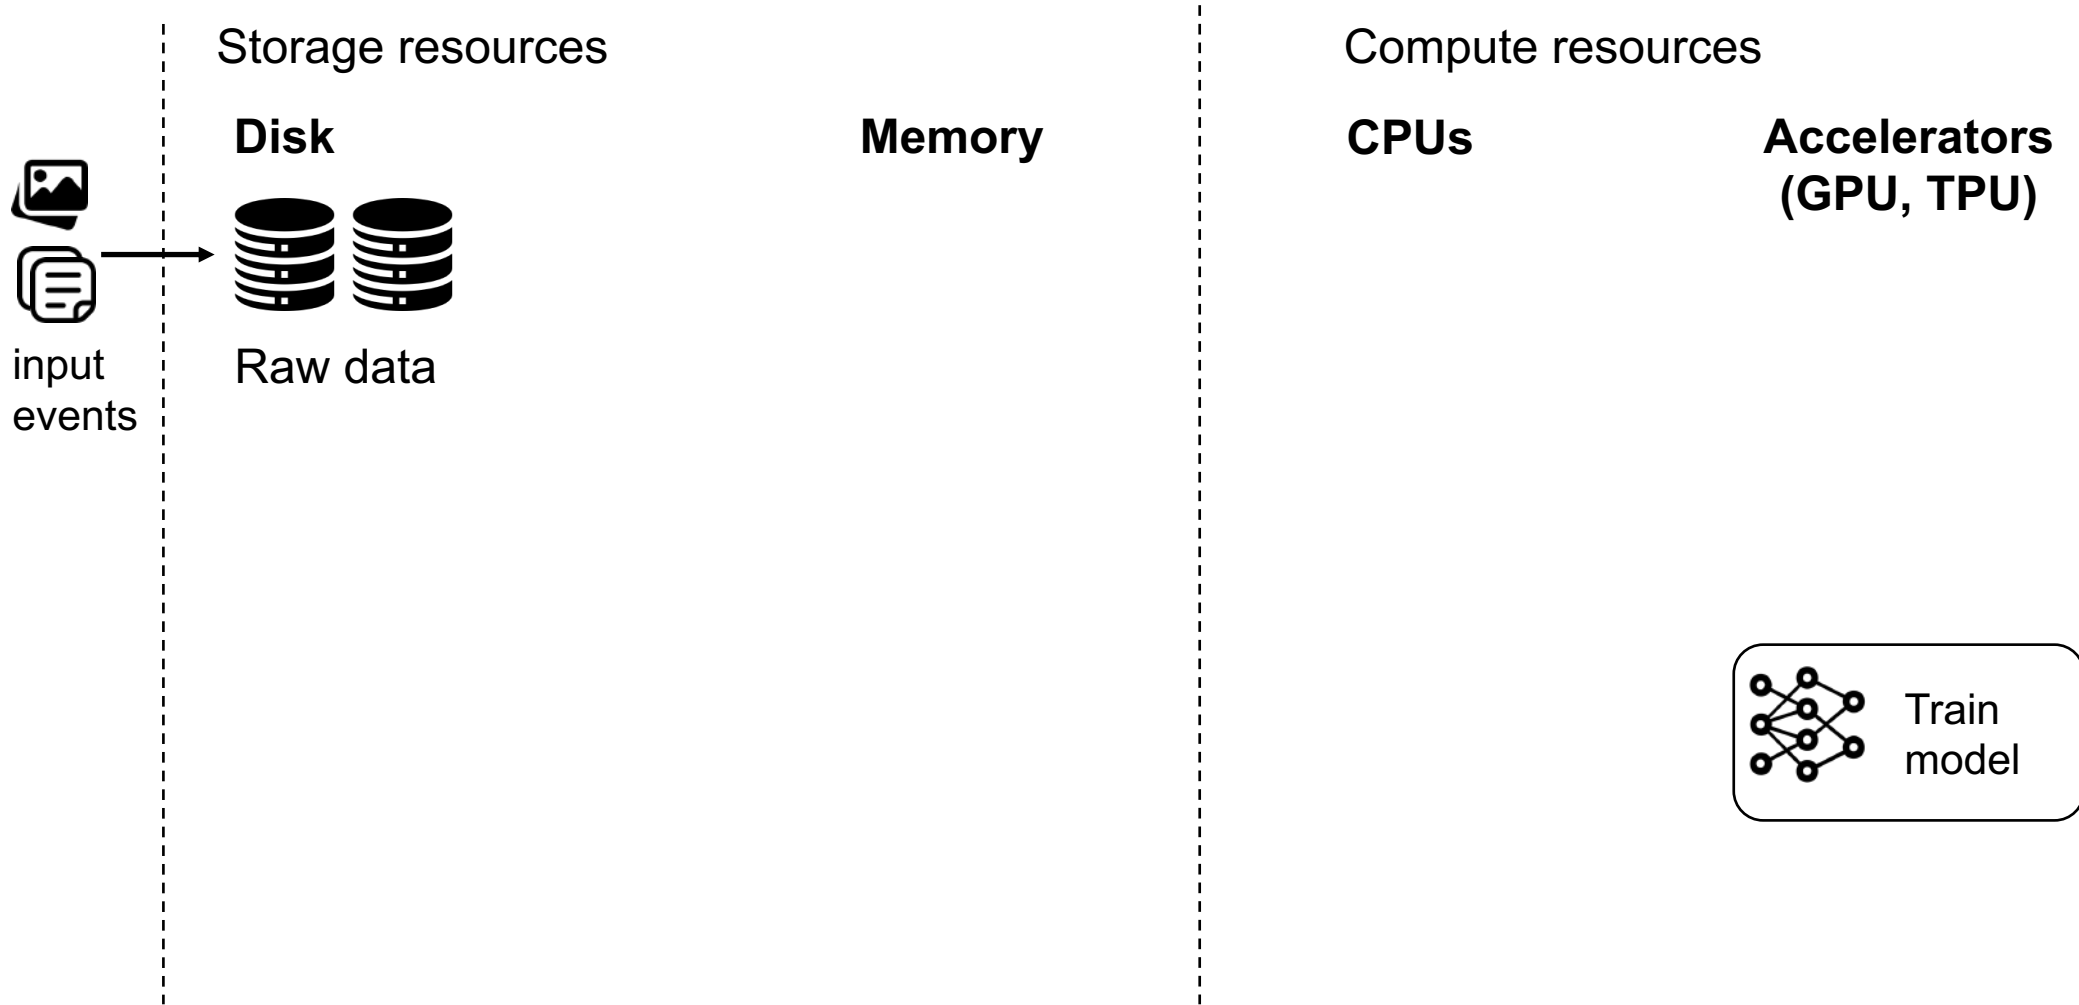
**Accelerators
(GPU, TPU)**



Train
model

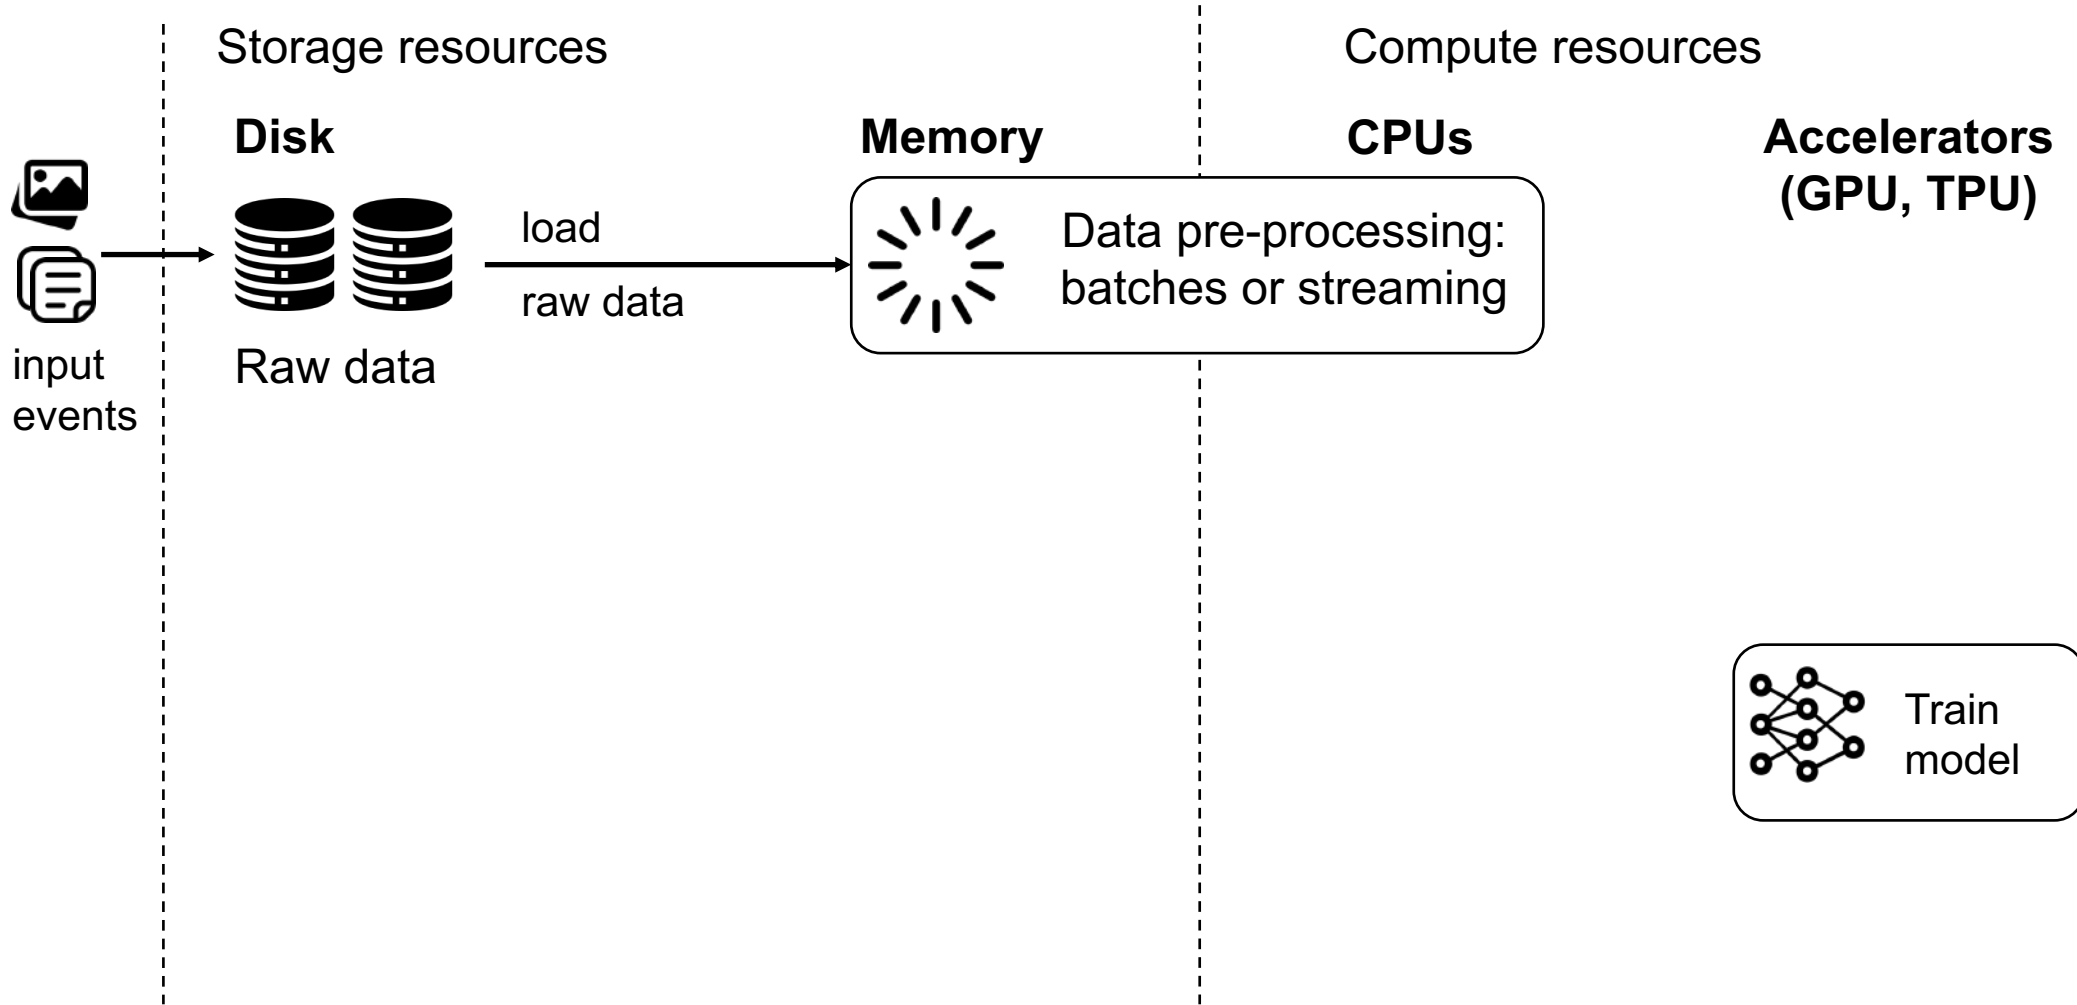


Data Pipeline in ML: Pre-processing



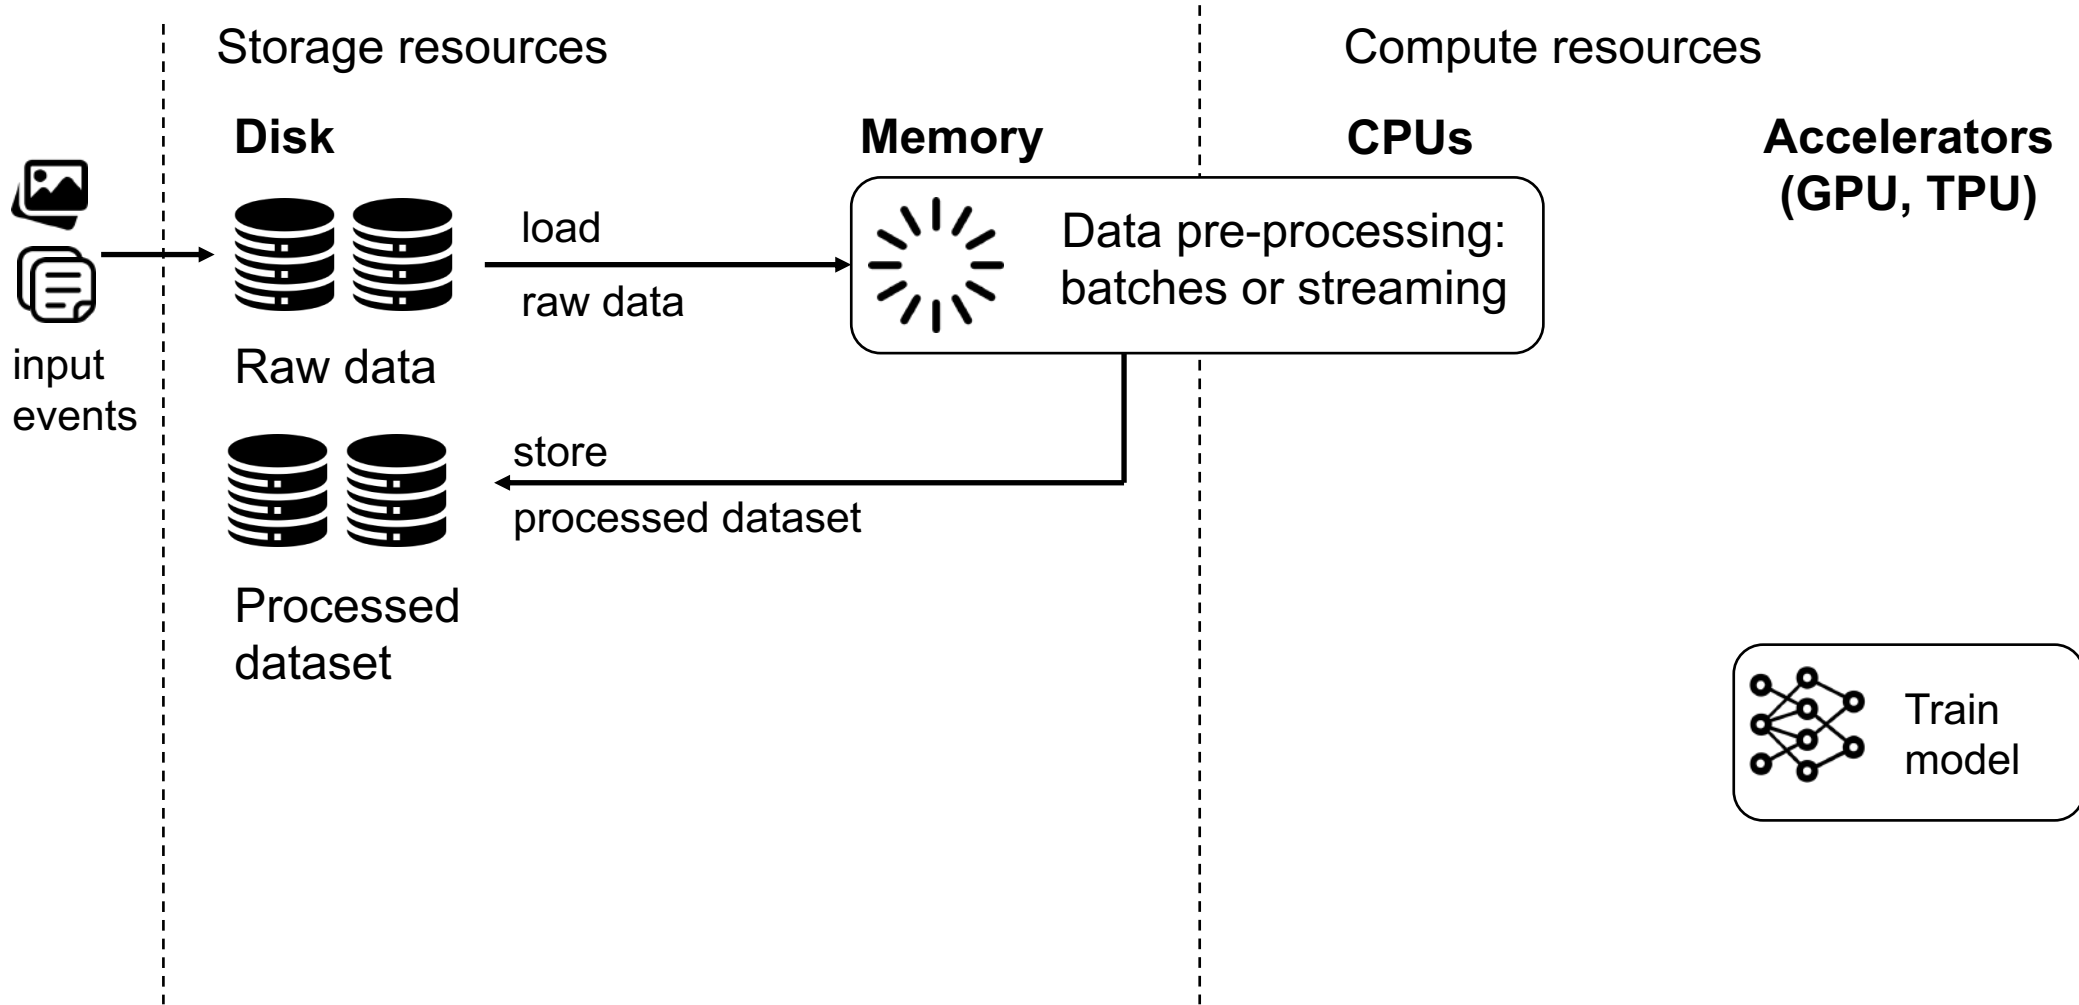


Data Pipeline in ML: Pre-processing






Data Pipeline in ML: Pre-processing




Stages of the ML Pipeline



[1] Murray et al. *tf.data: A Machine Learning Data Processing Framework*, VLDB 21.

[2] Zhao et al. *Understanding Data Storage and Ingestion for Large-Scale Deep Recommendation Model Training* ISCA 22.



Data pipeline in ML: Training

Storage resources

Disk




Cleaned
dataset

**System
Memory (DRAM)**

Compute resources

CPUs

**Accelerators
(GPU, ASIC)**



Data pipeline in ML: Training

Storage resources

Disk



Cleaned
dataset

 TensorFlow
PYTORCH

load
data


System
Memory (DRAM)

Cache
data

Compute resources

CPUs

Accelerators
(GPU, ASIC)



Data pipeline in ML: Training

Storage resources

Disk



Cleaned dataset

 TensorFlow
PYTORCH

load
data

System
Memory (DRAM)


Cache
data

Compute resources

CPUs

Online data
Pre-processing

Accelerators
(GPU, ASIC)



Data pipeline in ML: Training

Storage resources

Disk



Cleaned dataset

 TensorFlow
PYTORCH

load
data

System
Memory (DRAM)


Cache
data

Compute resources

CPUs

Online data
Pre-processing

Accelerators
(GPU, ASIC)




Load data
in batches

MLPerf Storage



Data cleaning &
pre-processing



Focus on **storage impact in ML/AI**


Realistic **storage** settings in
training phase

No accelerator required to run

Minimal AI/ML knowledge

MLPerf Storage – workloads

Workload	Image segmentation	Natural language processing	Recommender Systems
Model	3D U-Net	BERT	DLRM
Seed data	KiTS19 Set of images	Wikipedia 2020 Text	Criteo Terabyte Click logs
Framework	Pytorch	Tensorflow	Pytorch
I/O behavior	Random access inside many small files	Sequential access of small subset of files, streamed.	Random access inside one large file



<https://github.com/mlcommons/storage>

- **Single node**
- Many **simulated accelerators**.
- **Synthetic datasets** generated from real dataset seed.
- **Local storage**

MLPerf Storage – Benchmark metric

Must capture dynamics between storage and compute.

MLPerf Storage – Benchmark metric

Must capture dynamics between storage and compute.

Storage-centric metrics

- ✓ IOPS
- ✓ Latency
- ✓ Read/Write throughput
- ✓ Capacity

Compute-centric metrics

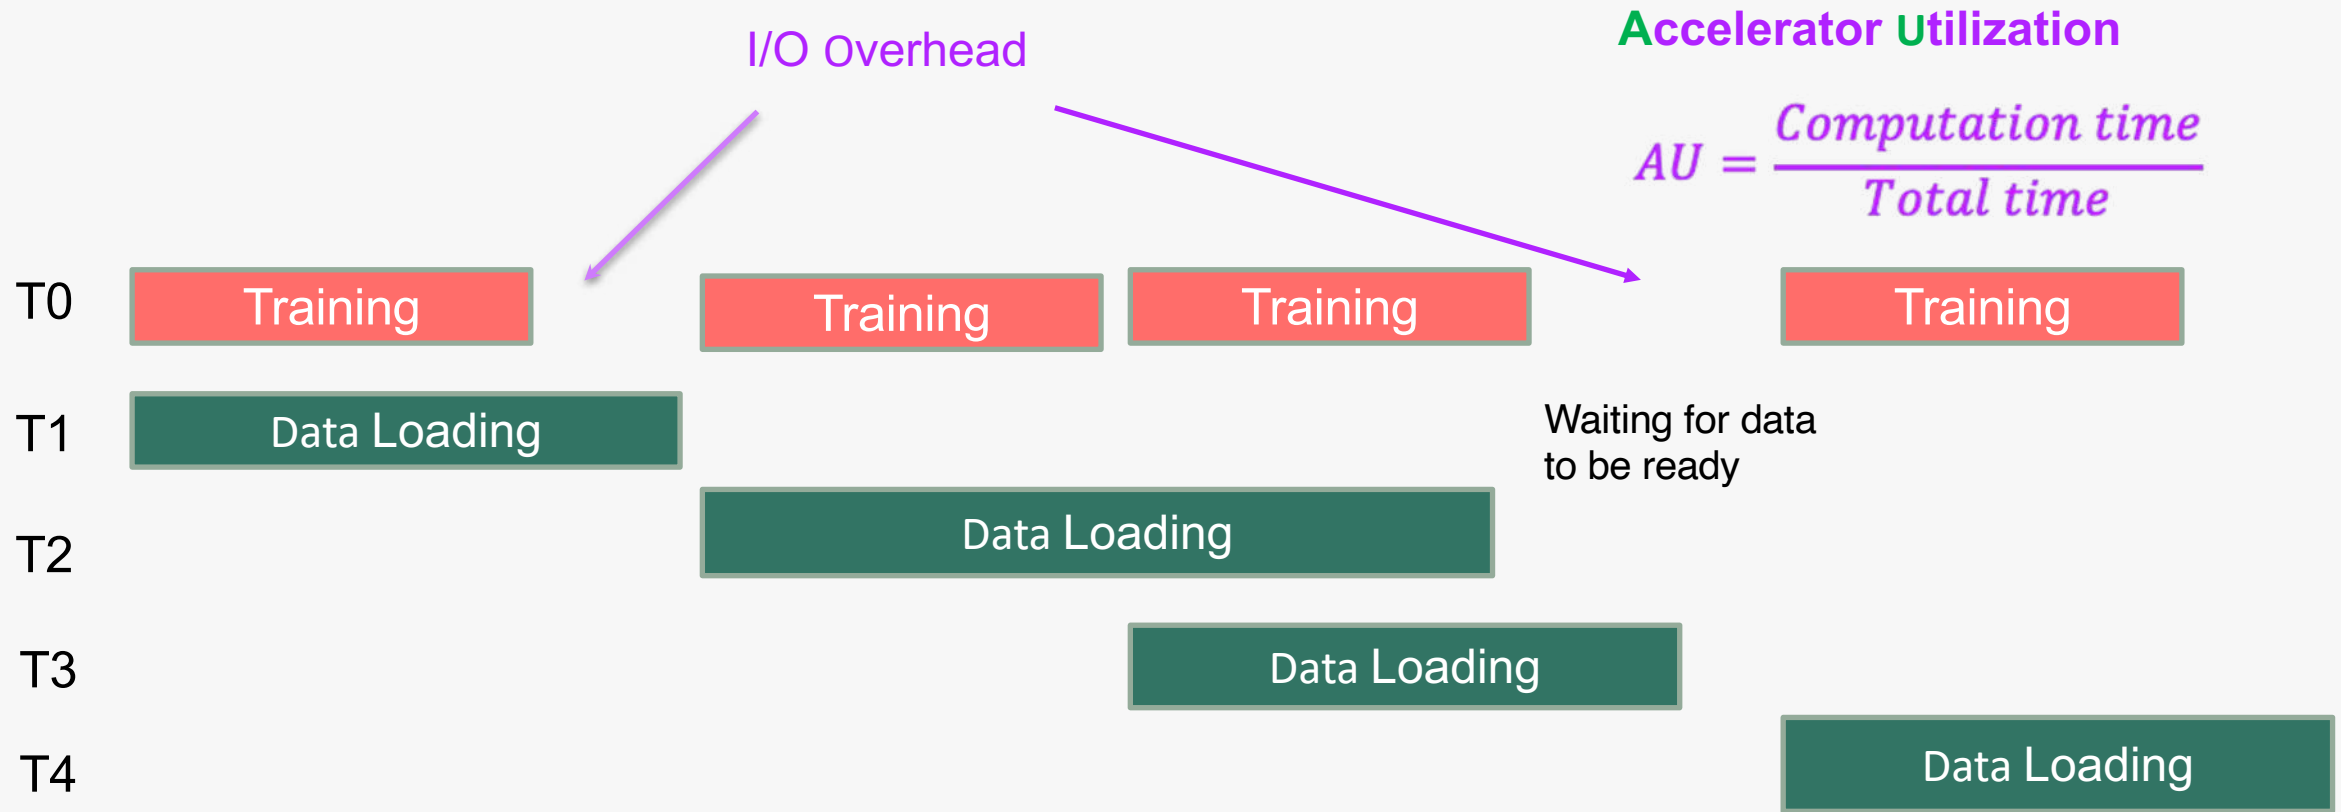
- ✓ Training time
- ✓ Trained model accuracy
- ✓ Accelerator utilization

☹️ Neither metric is enough to capture the storage-compute relationship

- Storage metrics too generic. Cannot capture dynamics of ML workloads.
- Compute-centric metrics too narrow (e.g., no notion of dataset size).

Proposed metric

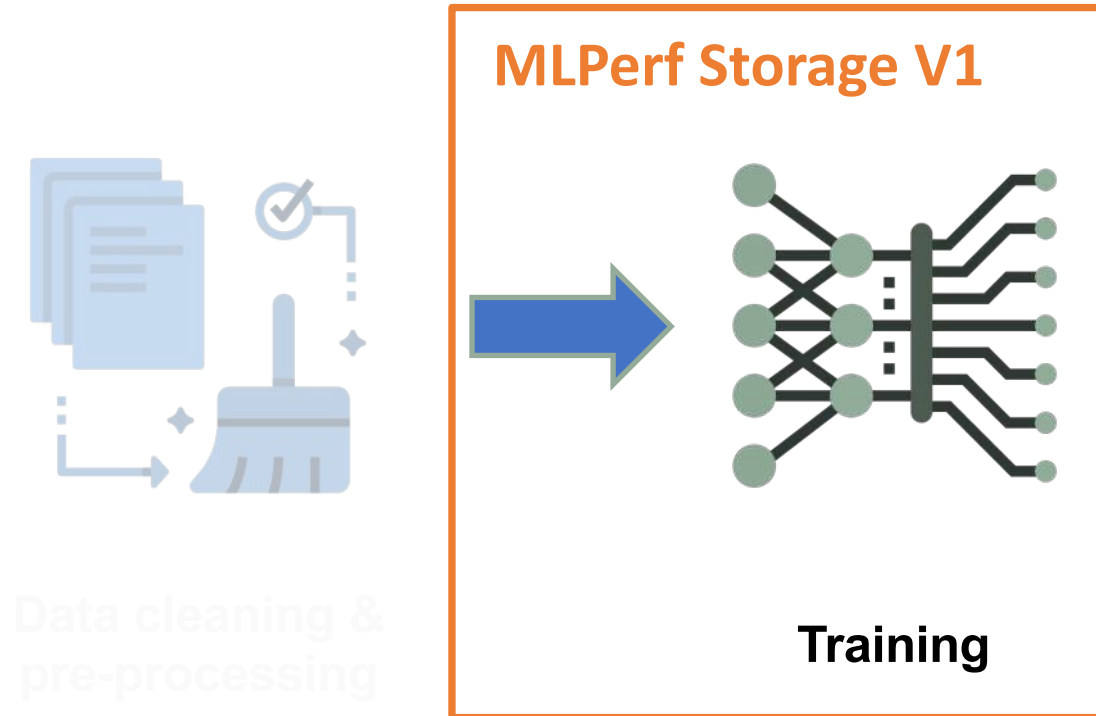
$$\text{Throughput} = \frac{\text{num_samples}}{\text{Time per epoch}}$$



Goal of benchmark:

Maximize samples / second, given an Accelerator Utilization > 90% at a certain scale.

MLPerf Storage




Focus on **storage impact in ML/AI**

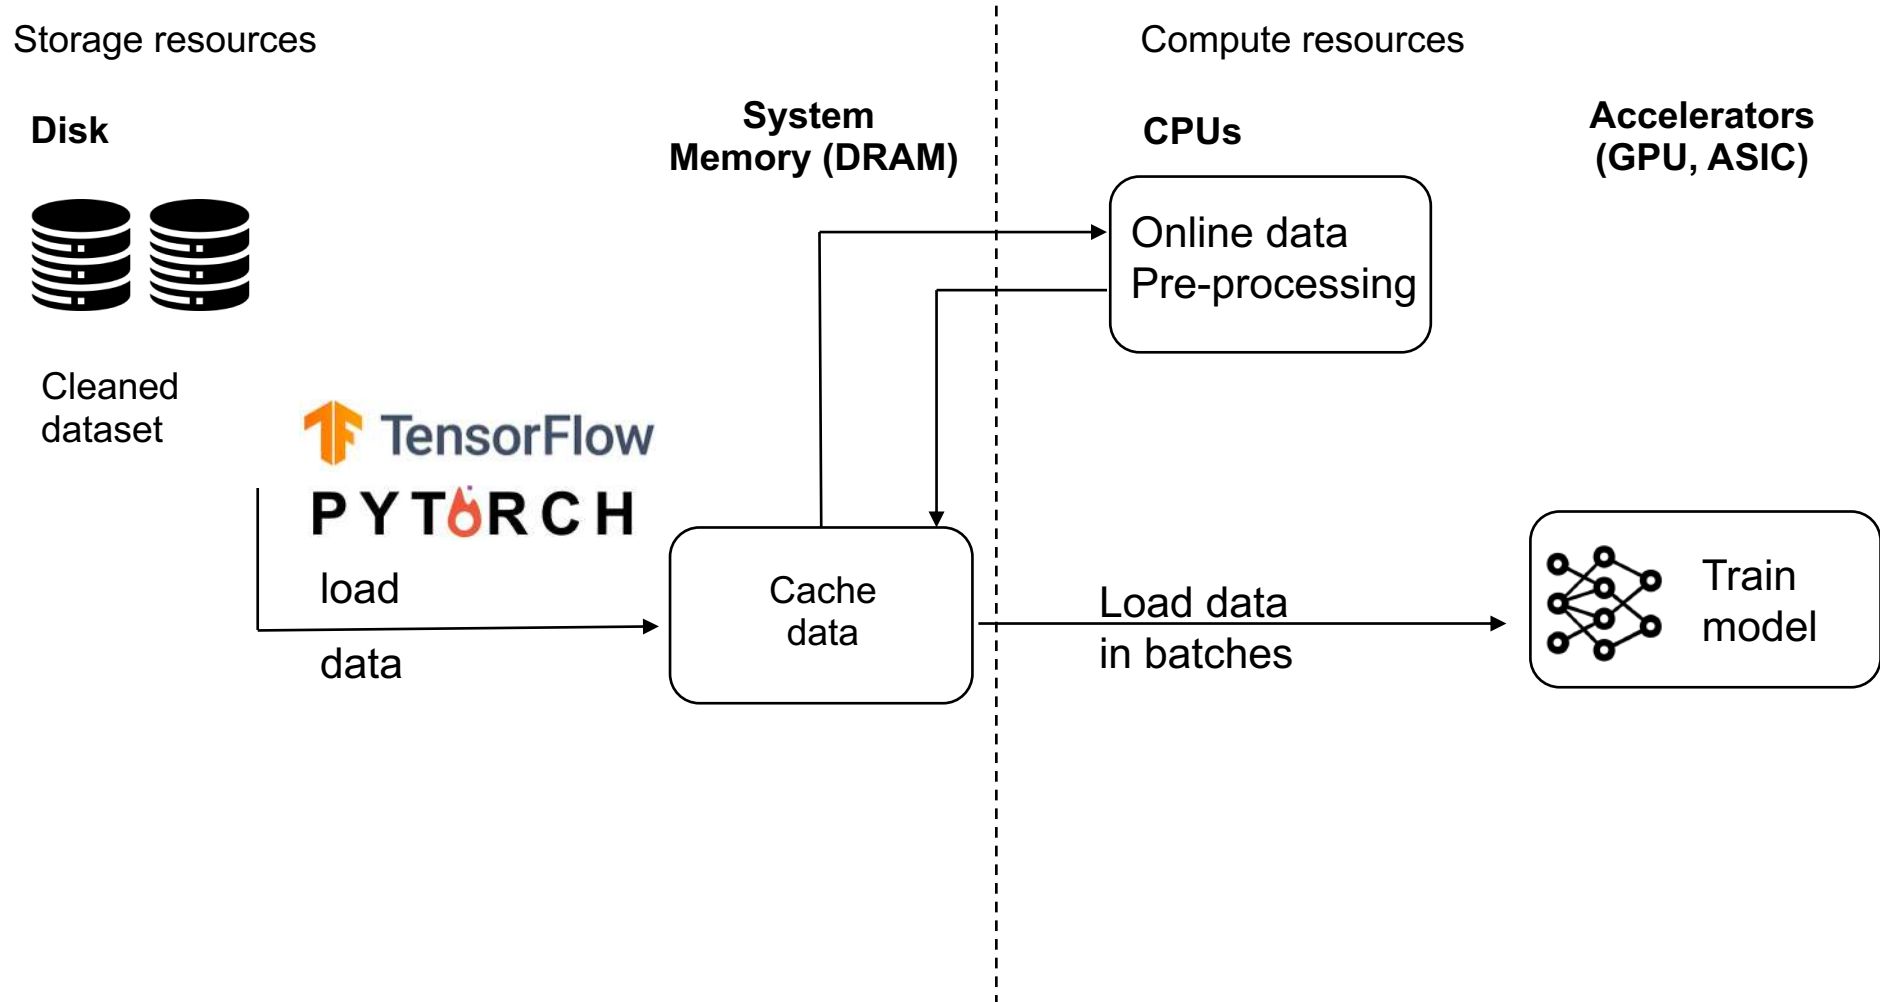
Realistic **storage** settings in
training phase


No accelerator required to run

Minimal AI/ML knowledge

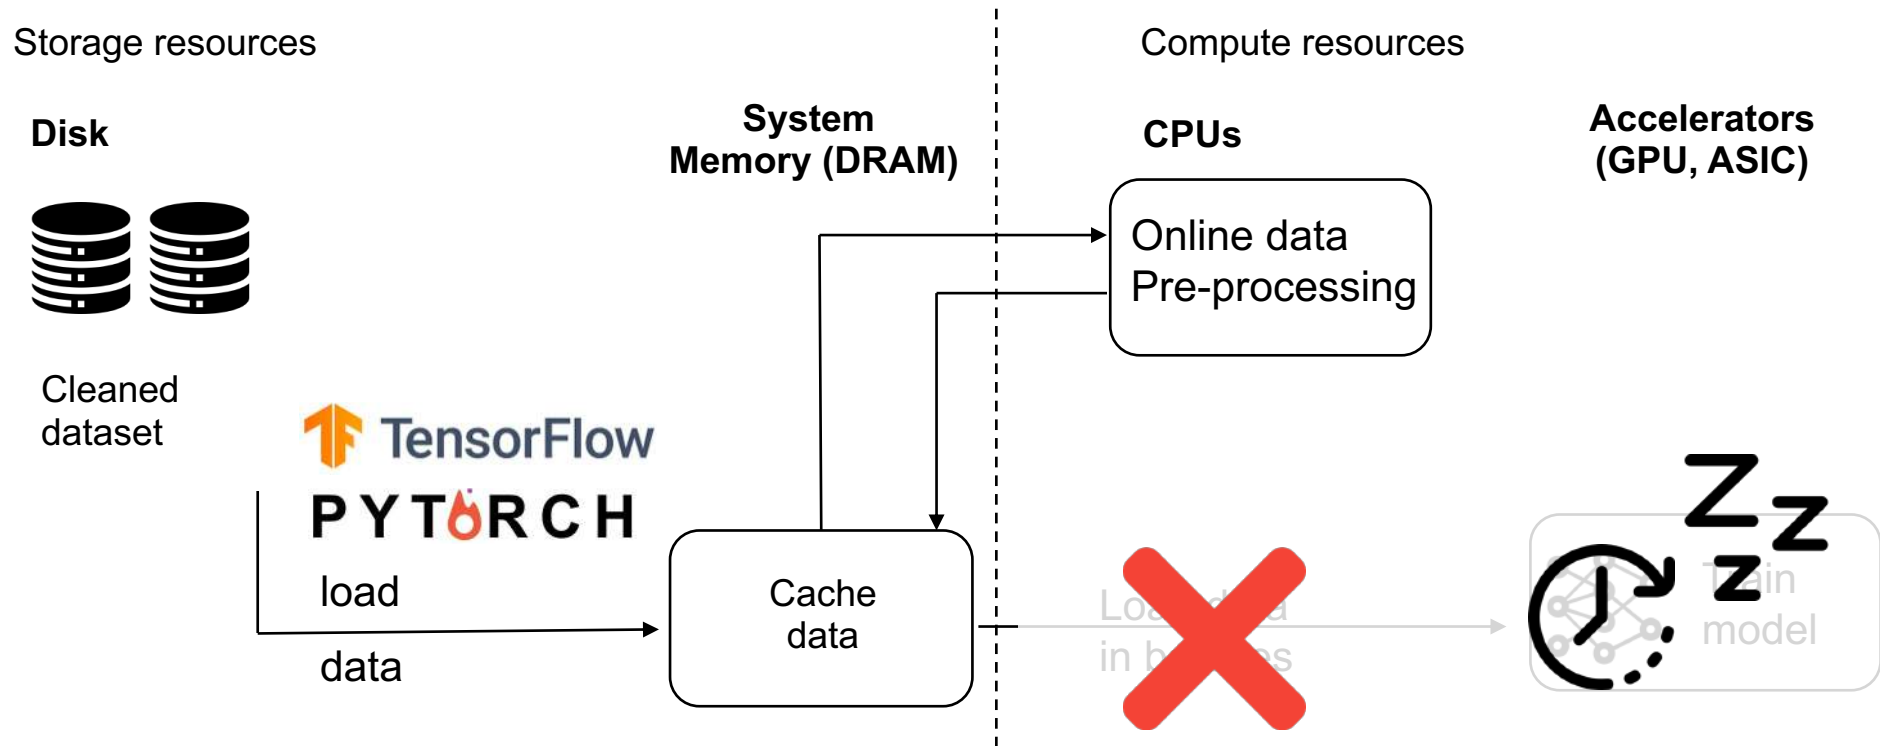


Data pipeline in ML: Training






Data pipeline in MLPerf Storage benchmark



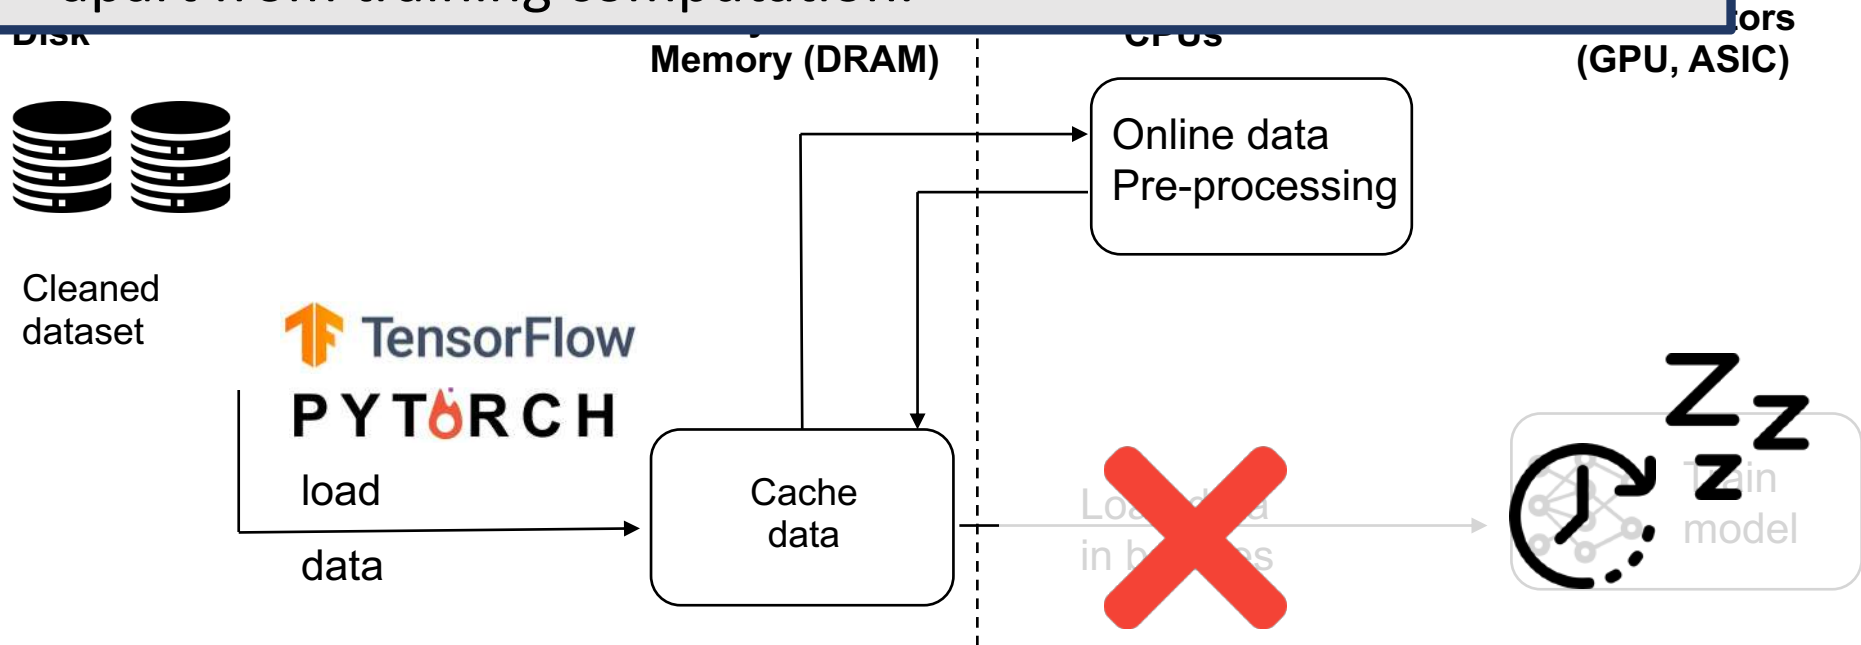
Benchmark is built as an extension of DLIO [1]

[1] H. Devarajan, H. Zheng, et al. DLIO: A Data-Centric Benchmark for Scientific Deep Learning Applications, CCGrid '21.




Data pipeline in MLPerf Storage benchmark

- ✓ Realistic storage settings: nothing changes in data pipeline, apart from training computation.



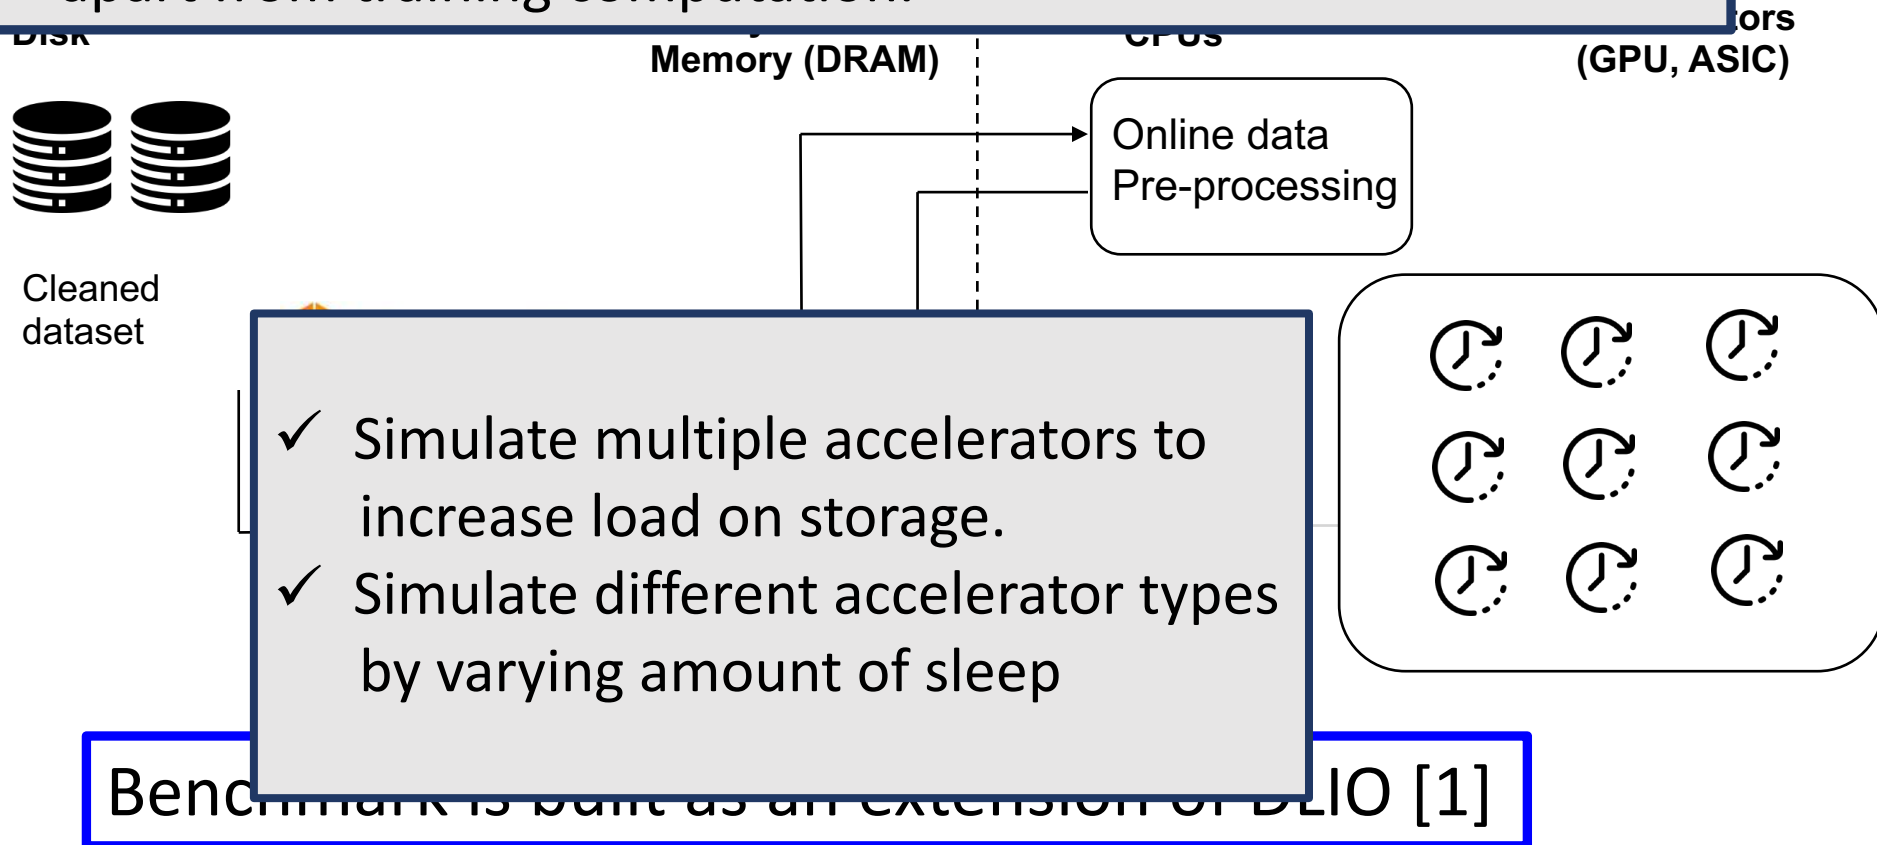
Benchmark is built as an extension of DLIO [1]

[1] H. Devarajan, H. Zheng, et al. DLIO: A Data-Centric Benchmark for Scientific Deep Learning Applications, CCGrid '21.



Data pipeline in **MLPerf Storage benchmark**

- ✓ Realistic storage settings: nothing changes in data pipeline, apart from training computation.



[1] H. Devarajan, H. Zheng, et al. DLIO: A Data-Centric Benchmark for Scientific Deep Learning Applications, CCGrid '21.

Experimental Evaluation


- DGX-1 server
 - 8 x V100 GPUs, 32GB GPU memory
 - 512GB DRAM
- Dataset size : Host memory size = 2:1

3D U-Net

- Pytorch, KiTS19 dataset seed
- Small model, large data.
 - 100s MB per sample
- One sample per file.

Simulating training time does not impact I/O patterns

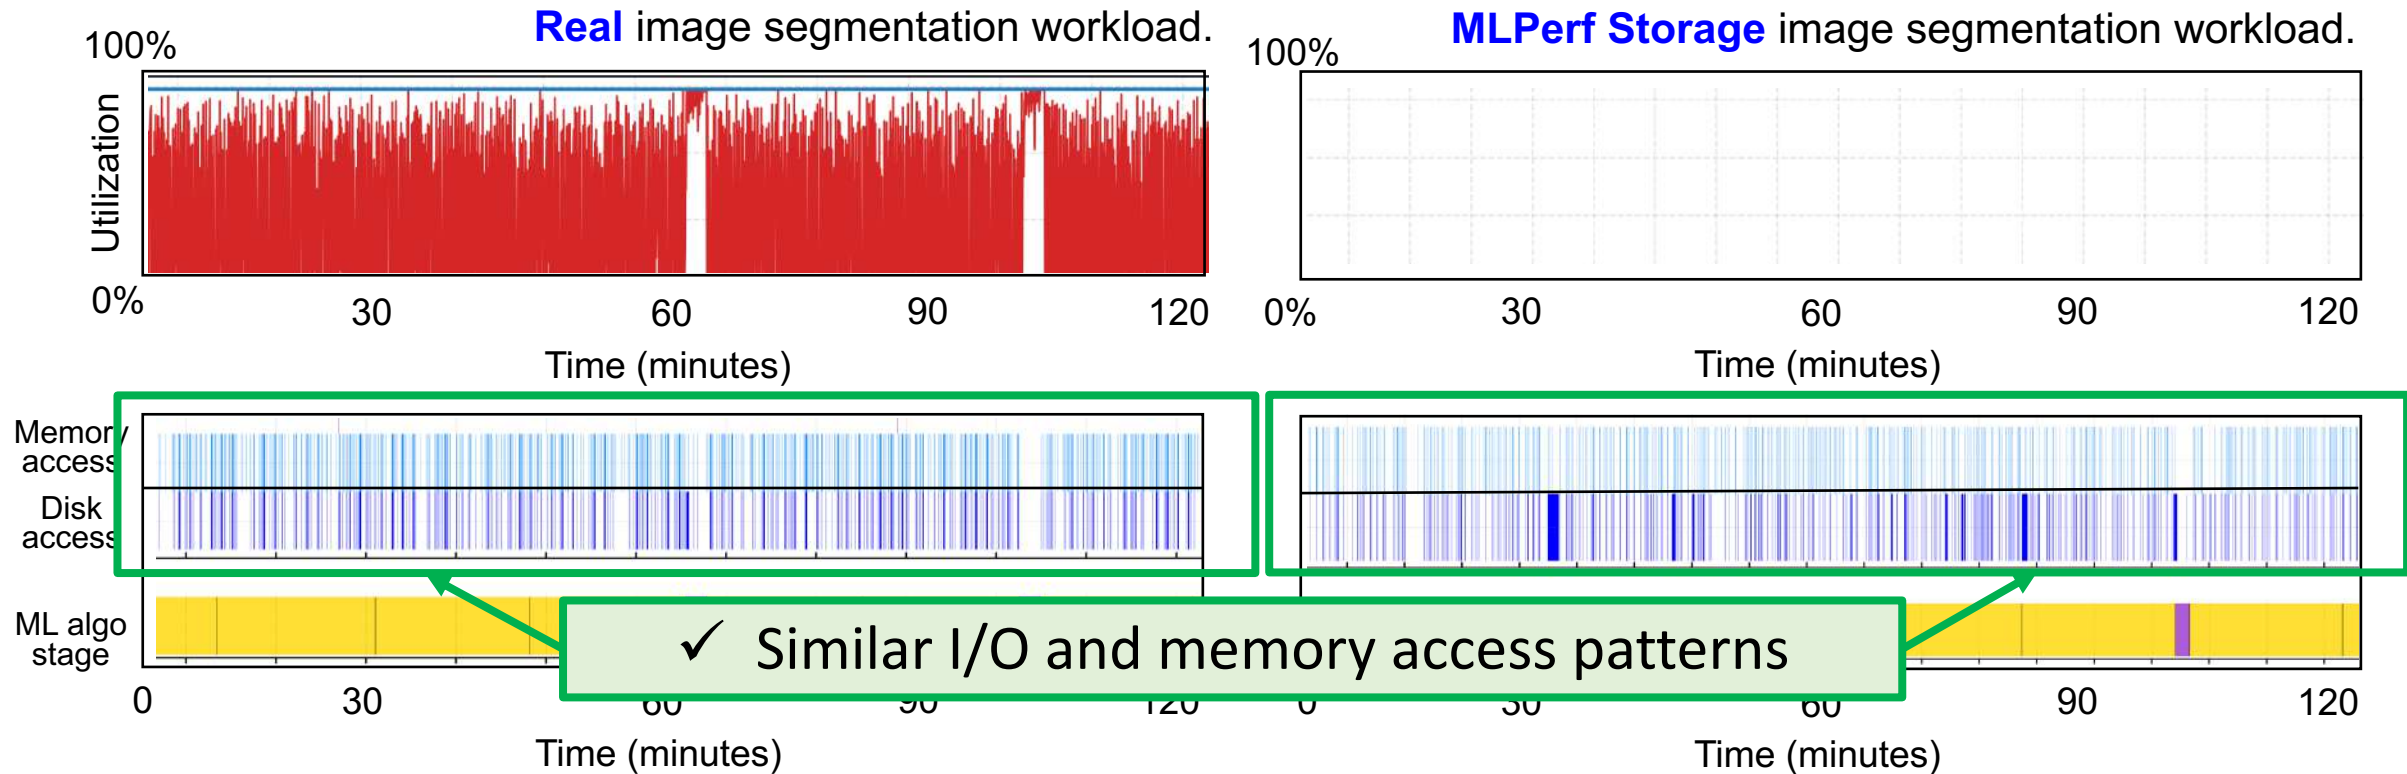
■ ML Training ■ ML Evaluation ■ Disk I/O Read ■ In-memory Read ■ GPU



Experiment setup: DGX-1 with 8xV100 GPUs, 512GB DRAM. Dataset : KiTS19, Dataset size:Memory size ratio 2:1

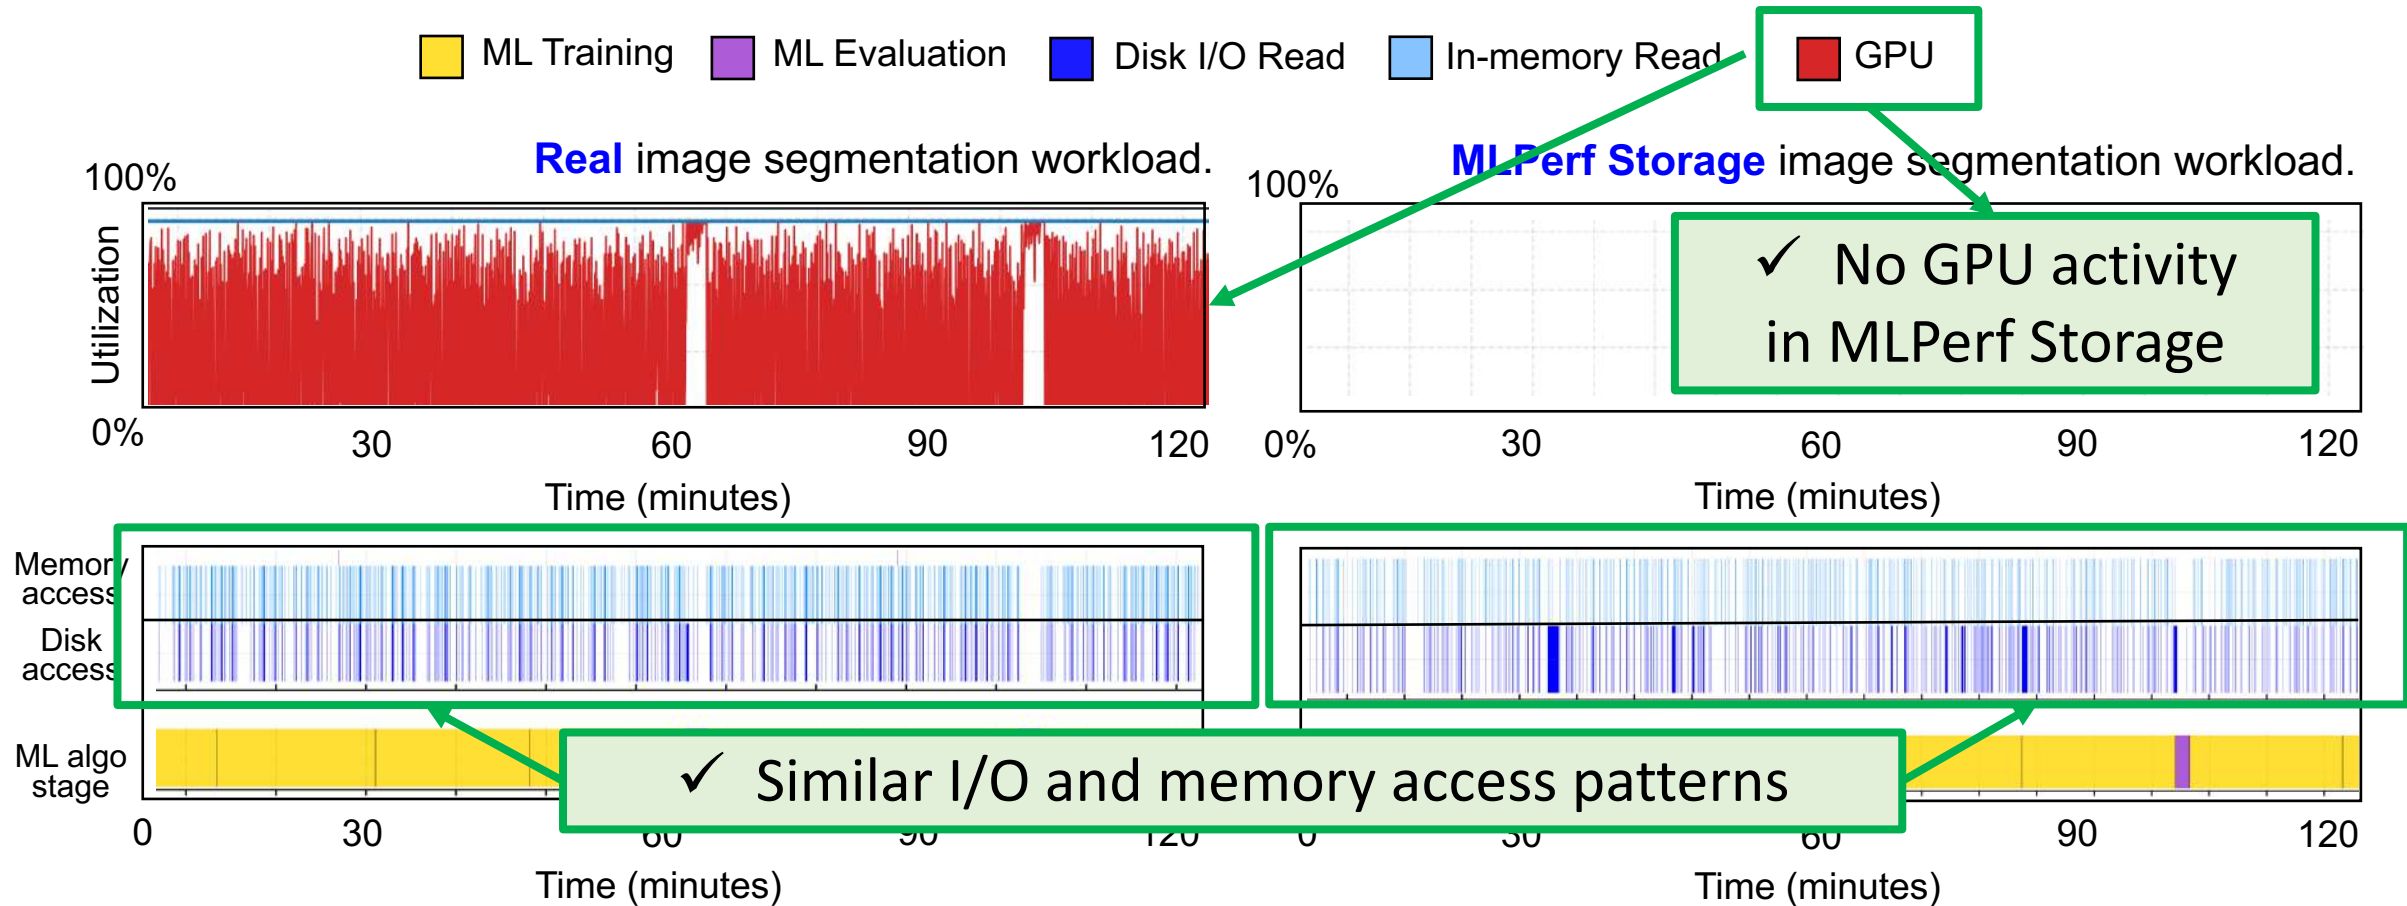
Simulating training time does not impact I/O patterns

■ ML Training ■ ML Evaluation ■ Disk I/O Read ■ In-memory Read ■ GPU




Experiment setup: DGX-1 with 8xV100 GPUs, 512GB DRAM. Dataset : KiTS19, Dataset size:Memory size ratio 2:1

Simulating training time does not impact I/O patterns



Experiment setup: DGX-1 with 8xV100 GPUs, 512GB DRAM. Dataset : KiTS19, Dataset size:Memory size ratio 2:1

3D U-Net: Step Breakdown




Natural Language Processing: BERT

- Tensorflow implementation, Wikipedia dataset seed
- Large model, small data.
 - Model takes up most of GPU memory
- Many small samples per file
 - ~300K samples per file
- Sequential access inside the files
 - Prefetching


BERT

ML Training ML Evaluation Checkpointing Disk I/O Read In-memory Read GPU




BERT

ML Training ML Evaluation Checkpointing Disk I/O Read In-memory Read GPU



Checkpointing
is a bottleneck
GPU → 0% utilization during checkpoint

BERT: Step Breakdown




DLRM


- PyTorch implementation, Criteo dataset seed
- Large model, Large data.
 - Model and data parallelism
- Many small random accesses inside a large file

DLRM

ML Training ML Evaluation Checkpointing Disk I/O Read In-memory Read GPU



DLRM: Step Breakdown



Lessons learned so far

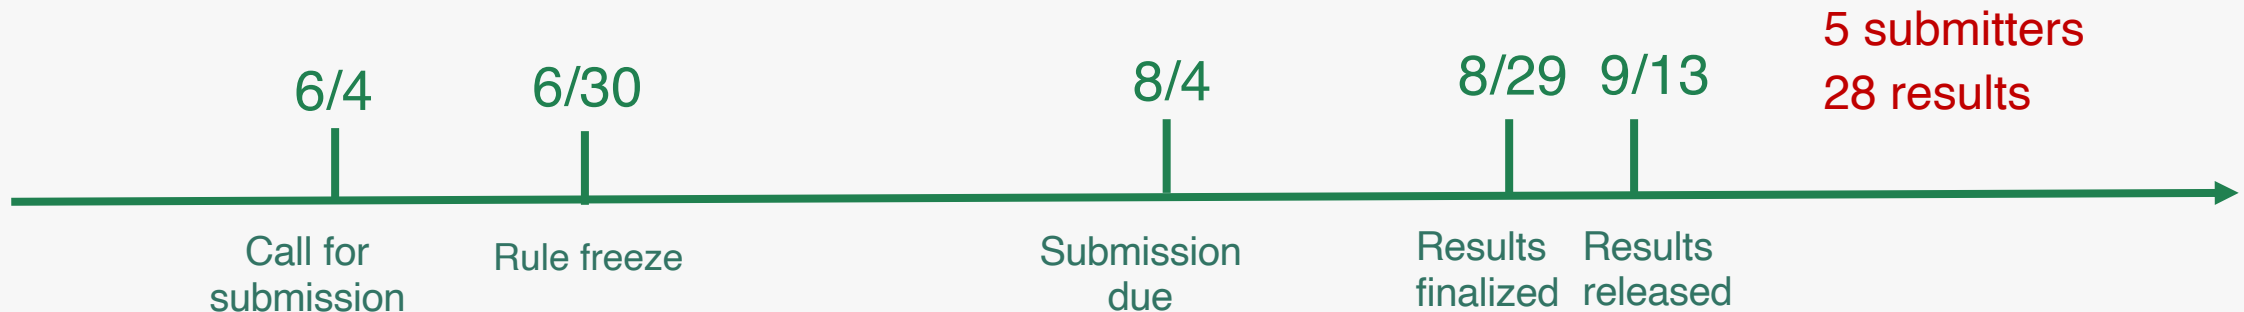
- Storage on its own is not the main issue.
- Studying I/O patterns revealed opportunities for improvement:
 - In the data loaders (3D U-Net)
 - In the checkpointing (BERT)
 - In the algorithm (DLRM)

MLPerf Storage v0.5 results overview



Workload selected

Workload	Image segmentation	Natural language processing
Model	UNet3D	BERT
Seed data	KiTS19 Set of images	Wikipedia 2020 Text
Sample size	~146 MB	~2.5 KB
Framework	Pytorch	Tensorflow
I/O behavior	Randomly select and read a file	Sequential access a subset of files, streamed

Timeline

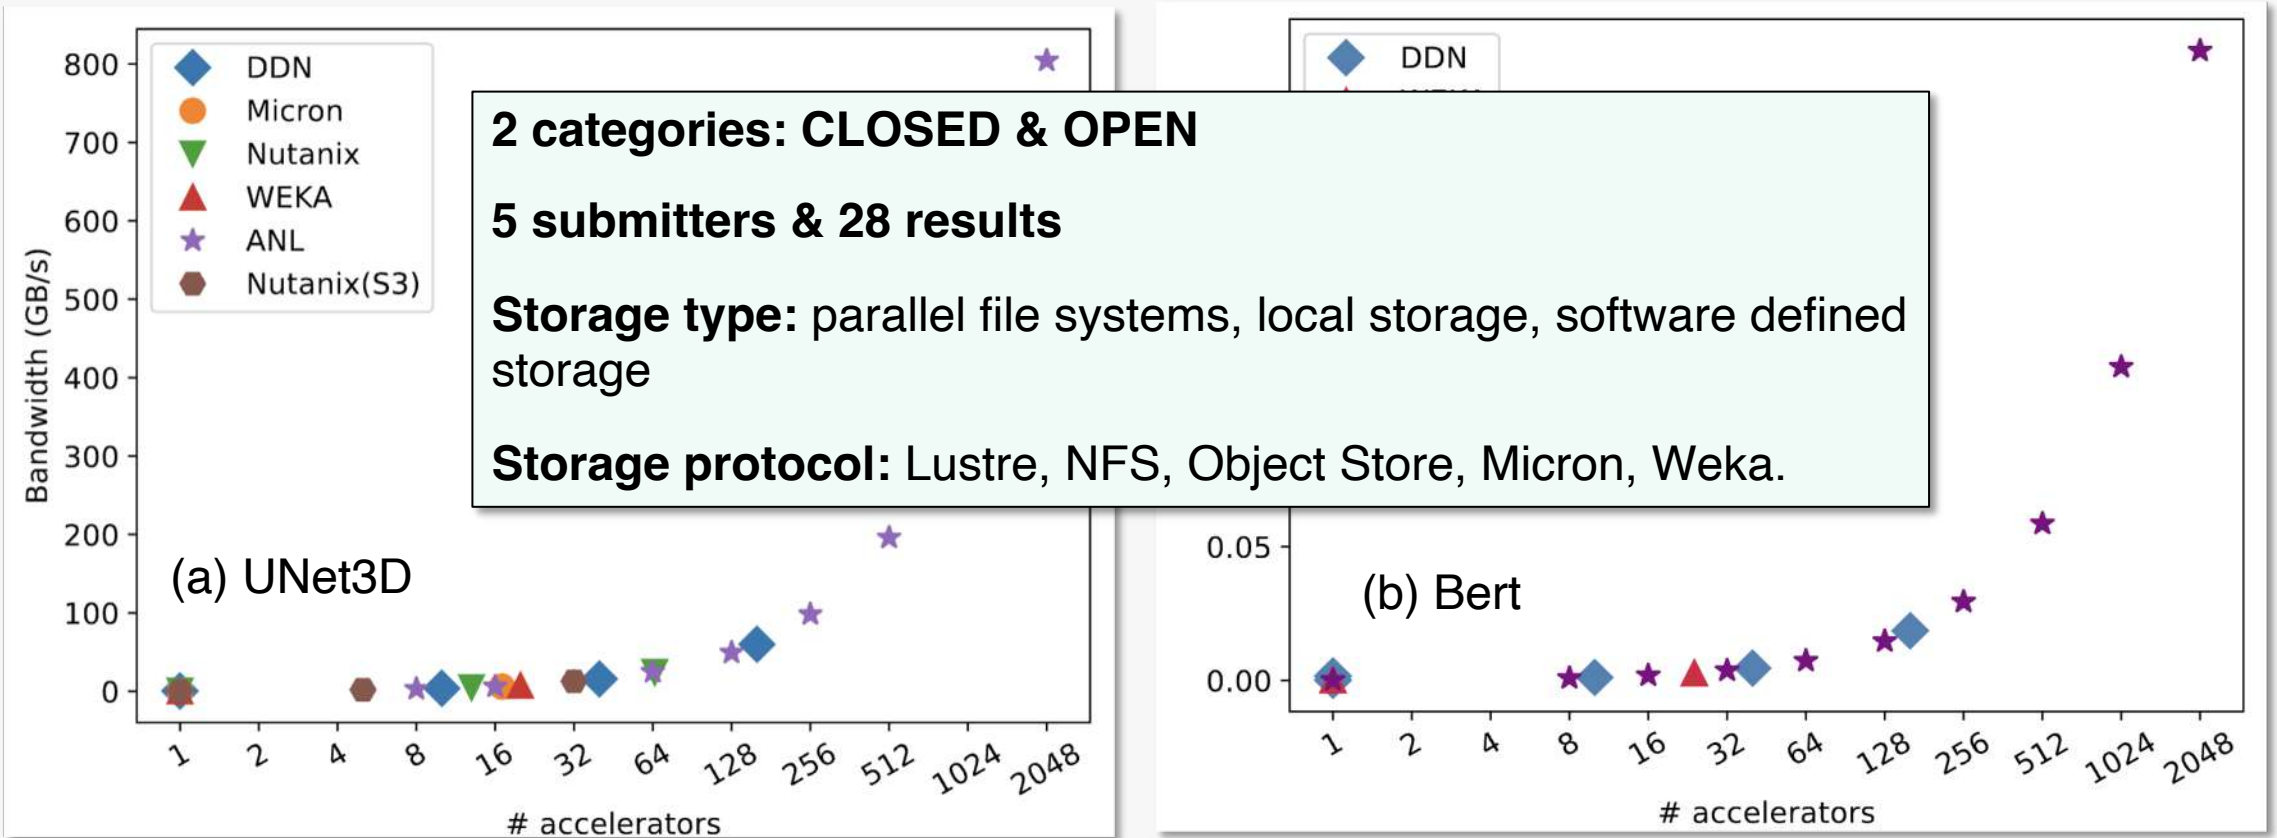


MLPerf Storage v0.5 results overview



Scatter plots of the results from the submitters: (a) UNet3D and (b) Bert. UNet3D is I/O intensive workload and Bert is compute intensive

MLPerf Storage v0.5 results overview



Scatter plots of the results from the submitters: (a) UNet3D and (b) Bert. UNet3D is I/O intensive workload and Bert is compute intensive

Next Steps in MLPerf Storage

Collect **processing times** for different accelerator types: A100, H100.

Benchmark competition round 2: <https://github.com/mlcommons/storage>

I/O in distributed training

New workloads (LLM, text-to-image, HPC)

Workload collocation

Extend benchmark with **ML pre-processing phase.**

McGill DISCS Lab



discslab.cs.mcgill.ca
gitlab.cs.mcgill.ca/discs-lab

Postdoctoral
Researcher



Dr. Stella Bitchebe

PhD
Candidates:



Nelson Bore



Jiaxuan Chen



Shubham Vashisth




Prithish Mishra



Rahma Nouaji



Masters
Students



Zachary Doucet



Aayush Kapur



Aidan Goldfarb



Ruoyu Deng

Key Takeaways – MLPerf Storage

MLPerf Storage is a new benchmark

Realistic **storage** settings

No accelerators required to run

Follow MLPerf Storage repository for updates:

<https://github.com/mlcommons/storage>

Get involved

<https://mlcommons.org/working-groups/benchmarks/storage/>

Share your thoughts
Email oana.balmau@mcgill.ca

Thanks to all working group co-chairs!



Curtis Anderson
Panasas



Huihuo Zheng
Argonne National Labs



Johnu George,
Nutanix